

AI Opportunity Roadmap

10 AI opportunities I'd actually prioritize if I joined Neeman's tomorrow, roughly in the order I'd tackle them. #1 is what I actually built for this assignment, so I've included it as the first entry rather than pretending it stands apart from the roadmap.

None of the hours-saved or KPI figures below are measured Neeman's results — they're my own estimates based on how big each problem looks from the outside. The only opportunity with real, working evidence behind it is #1, since I built and tested it myself.

P1 — I'd start this in the first 3 months

P2 — useful, after the first pilots land

P3 — valuable, but needs more data/readiness first

1

Business Analytics Copilot + Root Cause Analysis Agent

P1 — Harden & Productionize

BUSINESS PROBLEM

A likely challenge is that figuring out why revenue, category, channel, inventory, or marketing performance changed can require pulling numbers from multiple places and reconciling them manually. By the time there's a clean answer, the moment to act on it has often passed.

WHAT I'D BUILD

I built this for the assignment: a dashboard for what's happening, plus an agent that investigates why. The key design decision is that the AI never does the math itself a router first figures out what kind of question is being asked, plain pandas code calculates every number and evidence-strength label, and the AI's only job is to read that evidence and explain it in plain English. Every number it states gets checked against the underlying data before it's shown, so it can't just state something plausible-sounding that isn't actually true.

EXPECTED IMPACT

Should mean fewer “can someone pull me a number” requests to whoever owns reporting today, and gives teams like D2C, Marketplaces, and Marketing a way to dig into performance themselves instead of waiting in a queue.

Hours saved (est.): ~6-8 hrs/week for whoever currently does this reporting by hand

KPI: Time-to-insight; volume of ad-hoc analyst requests

Stack: Python, pandas for the math, LangGraph to sequence the agent's steps, Sarvam AI as the model, Streamlit for the interface

Why here: This is the one entry I can speak to with real confidence, since it's built and I've actually tested it everything else below is my best estimate from the outside.

BUSINESS PROBLEM

Neeman's is scaling from 46 to 100 stores. Every new store multiplies the number of SKU-store combinations that can quietly run low, and a stockout on a popular comfort SKU is a direct, avoidable loss.

WHAT I'D BUILD

A model that looks at how fast a SKU has been selling, how much stock is left, and typical resupply time, then flags SKU/store pairs likely to run out in the next couple of weeks as a ranked list for ops to review, not an auto-reorder. This is a forecasting problem, not a language problem, so I wouldn't reach for an LLM here at all; a fairly standard demand-forecasting model handles it well.

EXPECTED IMPACT

Should reduce lost sales from stockouts and cut the manual monitoring ops currently does by hand.

Hours saved (est.): ~8-10 hrs/week for the ops team, once tuned **KPI:** Stockout rate; inventory turnover

Stack: Python, pandas, a forecasting library (Prophet, or even a simpler exponential-smoothing model to start), Slack for alerts

Why here: High impact and reuses the data pipeline from #1 but I'd first confirm the inventory data actually carries reliable SKU-level lead times before committing real effort.

BUSINESS PROBLEM

Neeman's brand promise runs on comfort and consistently good reviews. But the complaints that don't fit that story — sizing issues, material problems, comfort regressions are scattered across the website, Amazon, Myntra, and support tickets, and nobody's really pulling those threads together today.

WHAT I'D BUILD

A pipeline that reads through reviews and tickets, groups similar complaints by SKU and material, and sends a weekly summary to product and QC. This is squarely an LLM use case understanding free-text complaints and clustering them by theme is exactly what language models are good at.

EXPECTED IMPACT

Faster detection of recurring product issues, which should feed back into product decisions and reduce returns tied to fixable problems.

Hours saved (est.): ~5-6 hrs/week for CX & product **KPI:** Return rate tied to sizing/material issues; CX resolution time

Stack: An LLM for classification, basic topic clustering, reusing the dashboard patterns from #1 for the digest

Why here: I'd want to check how much review data is actually exportable across all these platforms before committing that access is the real dependency, not the AI part.

BUSINESS PROBLEM

Marketing already tracks ROAS deterioration in the dashboard I built, but reacting to that shifting spend between channels is still manual and reactive.

WHAT I'D BUILD

Two separate problems here, and I'd treat them differently. Watching ROAS/attributed revenue by channel and flagging underperformers is a straightforward statistics/anomaly-detection problem, not an AI model. Where an LLM actually helps is drafting ad-copy variants for testing once a channel's been flagged.

EXPECTED IMPACT

Should shorten how long it takes to notice a channel going soft and actually react to it.

Hours saved (est.): ~6 hrs/week for marketing

KPI: Blended ROAS; time from problem detection to reallocation

Stack: Meta/Google Ads APIs, the existing marketing pipeline from #1, an LLM only for the copy-drafting piece

Why here: Useful, but depends on ad-platform API access I can't assume already exists that's a real integration project, not just a modeling task.

BUSINESS PROBLEM

Keeping product listings consistent across the website, Amazon, Myntra, and Flipkart for 150+ SKUs each with different formatting rules is tedious and easy to get wrong today.

WHAT I'D BUILD

Generate and maintain listing copy from one source-of-truth spec per SKU, respecting each platform's format rules. A clean LLM use case: writing on-brand copy within fixed constraints.

EXPECTED IMPACT

Faster, more consistent launches, and possibly a small conversion lift from tighter listings.

Hours saved (est.): This workload is genuinely periodic rather than steady I'd estimate ~2-4 hrs/week on average, concentrated around launch windows

KPI: Conversion rate on updated listings; launch cycle time

Stack: An LLM, marketplace seller APIs

Why here: Contained scope, easy to pilot on a handful of SKUs before rolling out wider.

BUSINESS PROBLEM

The analytics/RCA work above operates at the aggregate channel/category level. Individual store managers don't have an easy per-store view today, and that gap gets worse as the store count grows toward 100.

WHAT I'D BUILD

An automated daily or weekly digest per store (WhatsApp or Slack) sales, top/bottom SKUs, stock health, maybe a short AI-written note flagging anything unusual. Mostly a reporting/automation problem; the AI is a light summary layer on top of numbers pandas already computed.

EXPECTED IMPACT

Gives managers visibility without someone manually pulling numbers, and supports scaling store count without a proportional increase in oversight time.

Hours saved (est.): ~3-4 hrs/week per regional manager

KPI: Regional-manager oversight hours; time to catch a store-level issue

Stack: The same metrics pipeline as #1, extended with a `store_id` field, WhatsApp Business API or Slack

Why here: Honestly, this one's data-blocked right now the sales data I worked with doesn't carry `store_id`, so this can't really start until that's captured. Worth flagging early since it changes the sequencing.

BUSINESS PROBLEM

Footwear carries structurally high return rates industry-wide, and sizing is a plausible contributor to the return volume Neeman's processes.

WHAT I'D BUILD

A lightweight model estimating size-mismatch risk at checkout, using SKU and customer purchase history, triggering a fit-guide nudge before the order is placed. A standard supervised classification problem, not an LLM problem.

EXPECTED IMPACT

Fewer avoidable returns and lower reverse-logistics cost, protecting margin.

Hours saved (est.): ~5 hrs/week saved in reduced return-handling load

KPI: Return rate; reverse-logistics cost per order

Stack: A standard classifier (scikit-learn), checkout-flow integration

Why here: Needs labeled return/size data I haven't seen evidence of I'd treat this as needing real validation before committing serious time.

BUSINESS PROBLEM

A 250+ person, fast-growing, high-ownership team means new hires and cross-functional staff keep asking the same questions about SOPs, brand guidelines, and product specs.

WHAT I'D BUILD

A Slack chatbot that answers from internal documents. Technically this is “retrieval-augmented generation” a term that just means the system looks up the relevant document first and answers from that, rather than guessing from memory.

EXPECTED IMPACT

Faster onboarding and fewer repeat interruptions to managers.

Hours saved (est.): ~3-4 hrs/week per manager during onboarding waves

KPI: Onboarding time-to-productivity; repeat-question volume

Stack: An LLM plus a simple document search index, Slack bot

Why here: Low risk since it doesn't touch core business data, but also lower priority than the sales-facing work above a good “when there's spare capacity” project.

BUSINESS PROBLEM

Sustainable materials are core to the Neeman's brand story, but supplier compliance data is likely tracked manually today.

WHAT I'D BUILD

A tool that reads supplier compliance documents and tracks sustainable-material sourcing percentage over time, flagging gaps before they become a marketing-claim or investor-diligence problem.

EXPECTED IMPACT

Protects sustainability claims used in marketing, and produces cleaner metrics for investors given the company's stated path toward a public listing.

Hours saved (est.): ~3 hrs/week time

KPI: % sustainable sourcing tracked/compliant; audit-prep

Stack: Document extraction tools for reading supplier PDFs, the same reporting stack as #1

Why here: Lower urgency than the sales-facing opportunities, and depends on getting supplier documents into a consistent digital format first.

1
0

Automated Investor & Leadership Reporting

P3

BUSINESS PROBLEM

As a Series B company with a stated eye toward a public listing, board decks and investor updates are probably assembled by hand today with real risk that the numbers in them quietly drift from what the operational dashboards actually show.

WHAT I'D BUILD

A monthly deck generator pulling from the same underlying numbers as the internal dashboard, with AI drafting the narrative sections under the same never-invent-a-number discipline the RCA agent uses.

EXPECTED IMPACT

Saves prep time, and more importantly makes it much harder for investor-facing numbers to silently diverge from what operations actually sees.

Hours saved (est.): This is monthly work, roughly ~2 hrs/week averaged out

KPI: Deck-prep hours; manual reconciliations avoided

Stack: python-pptx, the same metrics pipeline, an LLM for narrative text with the same grounding check as #1

Why here: A genuinely good audit-readiness win, but I'd sequence it after the sales-facing pilots it touches leadership-visible material, so it needs to already be trustworthy on day one, not proven through iteration.

Prioritization & ROI

Not all ten of these matter equally right now. The table below is what I'd actually use to sequence work — not the numbered list order above. Impact/Effort/Data Readiness are directional High/Medium/Low calls meant to support a prioritization conversation, not precise scores.

#	Opportunity	Type	Impact	Effort	Data Readiness Today	Time to Value
1	Business Analytics Copilot + RCA	Foundation	High	Delivered	High	Immediate
2	Inventory & Stockout Prevention	Quick Win	High	Medium	Medium	Weeks
3	Customer Review & Feedback Intelligence	Quick Win	Medium-High	Medium	Low-Medium	Weeks-Months
4	Marketing Spend & ROAS Monitoring	Strategic Bet	High	Medium-High	Medium (needs ad API access)	Months
5	Marketplace Listing Copilot	Quick Win	Medium	Low-Medium	Medium	Weeks
6	Store Performance Digest	Quick Win (data-blocked)	Medium	Low	Low (no store_id today)	Months, after data capture
7	Return & Size-Fit Risk Prediction	Strategic Bet	Medium-High	High	Low	Months-Quarters
8	Internal Knowledge Copilot	Foundation	Medium	Medium	Medium	Months
9	Sustainability & Material Sourcing Tracker	Strategic Bet	Medium (risk mitigation)	Medium	Low	Quarters
10	Automated Investor & Leadership Reporting	Foundation	Medium-High (audit value)	Medium	High	Months

Measurement Philosophy

Before I'd call any of these “live,” I'd want three things nailed down: a baseline KPI measured under today's manual process, a target KPI the AI system is expected to move, and a defined post-launch measurement window (4–6 weeks feels about right) to check whether the target was actually hit before scaling further. The RCA agent from #1 already applies a version of this discipline at the level of a single answer every number it states gets checked against the evidence before it's shown. I'd hold the same standard at the initiative level: don't call something a win without a before/after number to back it up.

Realistic 30/60/90-Day Plan

This assumes one person with normal access to stakeholders and data owners not a team, and not simultaneous work on all ten opportunities. The goal is to ship one thing properly rather than prototype ten things badly.

Days 1–30 — Harden what's delivered, scope what's next

- Get real analyst/ops feedback on the Business Analytics Copilot + RCA Agent (#1) from actual use, not just assignment review and fix whatever breaks in practice.
- Check data readiness for the two most promising next opportunities: does inventory data actually carry SKU-level lead time and sales-velocity history for #2? Does the sales dataset capture store_id yet, which #6 needs?

- Pick one opportunity to take into a real pilot most likely #2, given its impact and the fact it reuses the pipeline I already built.

Days 31–60 — Pilot one opportunity for real

- Build and pilot #2 against a subset of stores/SKUs, using real historical sales and inventory data.
- Record the baseline stockout rate and inventory turnover for that subset before the pilot starts — that's the number everything else gets judged against.
- In parallel, at lower effort, scope a minimal version of #3 (Review Intelligence) against a single review source — a full multi-source pipeline is a Day-90+ problem, not a Day-60 one.

Days 61–90 — Measure, decide, plan next quarter

- Compare the pilot's stockout rate/turnover against the Day-30 baseline. Go/no-go based on that number, not on how the pilot felt.
- Document the data-readiness gaps (store_id capture, ad-platform API access, consolidated review data) blocking #4, #6, and #7 these become next quarter's actual prerequisite work, not a vague “later.”
- Propose next quarter's single next opportunity to stakeholders, backed by the Day 60–90 pilot results not a commitment to build all ten.

Prepared as Part 3 of the Neeman's AI Intern assignment.